

Redis Semantic Router for Document Classification

Wayne Cheng

User Story and Pain Points

- ~100K news articles per day need to be classified into 1 of 5 topics: business, entertainment, politics, sport, and tech
- Using a LLM (GPT) to classify the article
- Pain points:
 - Latency → bad user experience
 - Cost
 - Accuracy

Possible Solution

- Use a semantic router to classify articles
 - A semantic router takes the "semantics" (meaning) of text, and returns a "route" (class)
 - It's built with a vector database that returns the closest matching "route"
- Compared with a LLM, a semantic router is faster, less expensive, and possibly more accurate

business

article 1 embedding

article 2 embedding

⋮

article n embedding

sports

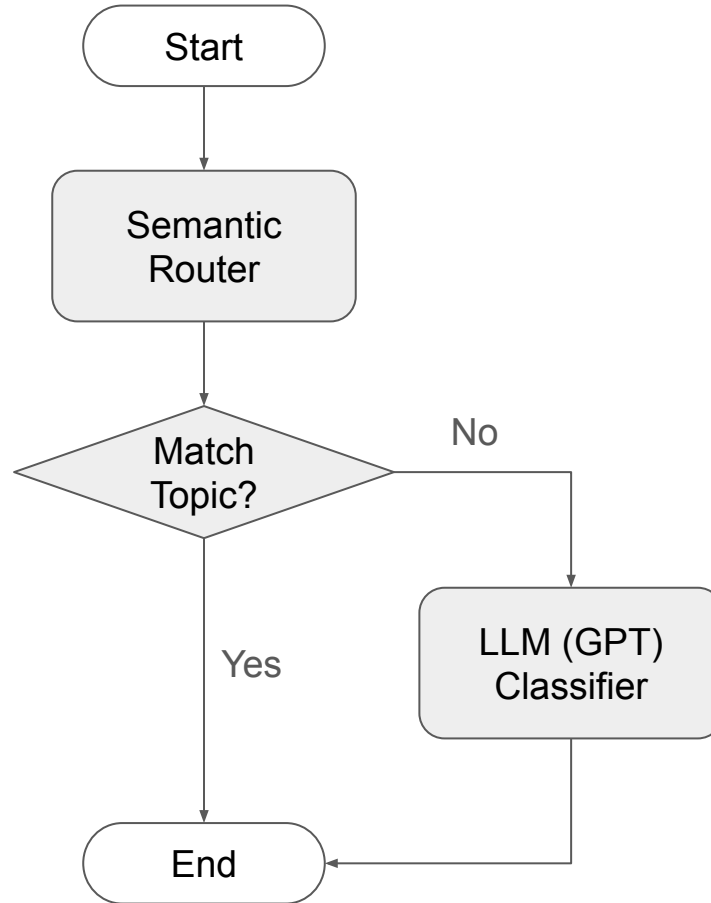
article 1 embedding

article 2 embedding

⋮

article n embedding

Architecture



Proof-Of-Concept Demo

- Github repository:
 - <https://github.com/audoir/redis-project>
- Prerequisites: Docker and Python must be installed
- Follow the instructions in the README.md file
- The configurations can be modified in src/config.py

Dataset

- The dataset contains articles, each labelled with a topic
- It is split into a test and train dataset
- Test dataset: 10 articles per topic (50 total)
- Train dataset: 50 articles per topic (250 total)

Baseline

- The test dataset is classified using a LLM (gpt-5-nano)
 - Simple classification task only requires a small model
 - Cost for input: \$0.05 per million tokens
 - Cost for output: \$0.40 per million tokens
- Average Latency: 5 seconds
- Average Cost = \$0.00025
- Accuracy = 88%

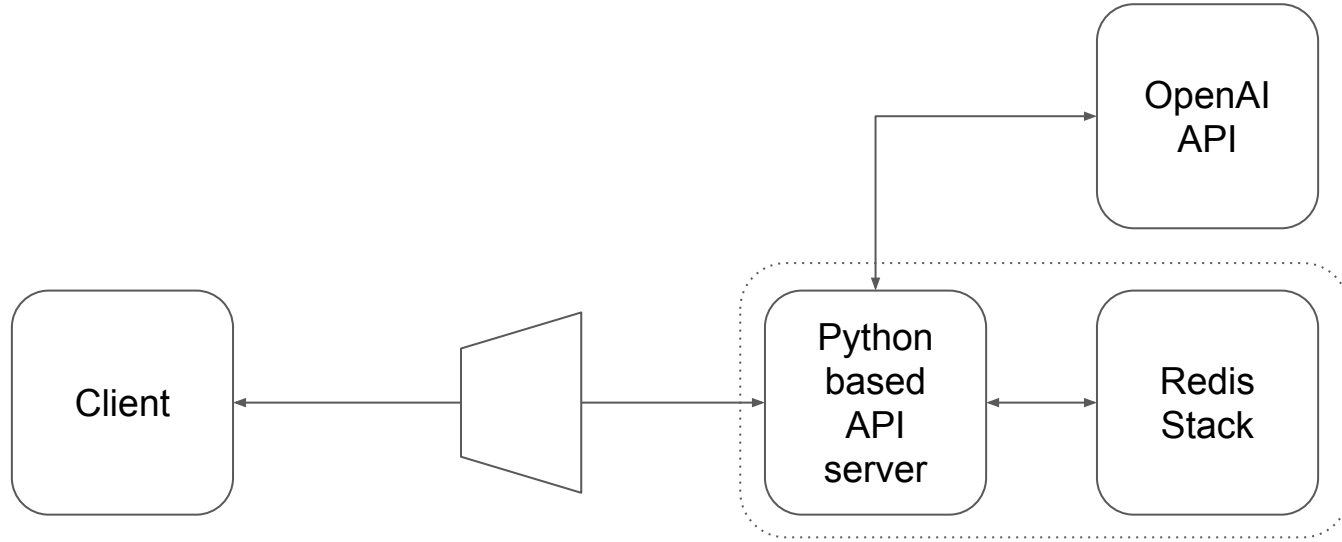
With Semantic Router

- The train dataset is used for the routes
 - Hugging Face's Sentence Transformers model (all-mpnet-base-v2) is used as the vectorizer
- The router distance threshold is tuned to 0.6 as a good balance between non-matches and false matches
 - The lower the threshold, the more non-matches
 - The higher the threshold, the more false matches
- Average Latency: 1 second
- Average Cost = \$0.000034
- Accuracy = 92%
- Router match = 88%

Comparison


- Percentage Improvements:
 - Latency: -79%
 - Cost: -86%
 - Accuracy: 4.6%
- For 100K articles:
 - Latency: **-112 hours**
 - Cost: **-\$22** (will more than cover the cost of running Redis for a month)
 - Correct: **4000 articles**
- The semantic router approach **reduces latency and cost**, while **improving accuracy**

System Design with Provisioned Compute



API server and Redis server can be replicated as needed based on demand

System Design with Serverless Services



Simple design → fast setup and auto-scaling

Extensibility

- Proactively create newer versions of the router with recent articles to maintain high accuracy, or add more topics

Why Redis?

- High performance using in-memory store
- Mature and well-adopted technology with lots of support

Fast and **easy** setup:

- RedisVL → build a semantic router with fewer lines of code
- Redis Cloud → fast prototyping and deployment; free to try

Implementation Strategy

- If you don't currently have Redis in your system, use Redis Cloud
 - Build a prototype API server using the demo code as a template
 - Duplicate the existing traffic to the prototype, and tune the configurations
 - Once the latency and accuracy is satisfactory, switch over to using this architecture for production
-
- The total implementation time is a few hours

Recap

- Using a semantic router for classification task improves the user experience by **reducing latency** and **improving accuracy**, while **reducing cost**
- You may have to tune with train dataset size, distance threshold, and vectorizer to get the best results
- You can either deploy Redis in your own system, or use Redis cloud for fast setup

Additional Resources

- Redis Vector Library: <https://docs.redisvl.com/en/latest/>
 - AI-native Python client library used in this demo
- Redis Cloud: <https://redis.io/docs/latest/operate/rc/>
 - Fully managed Redis database on the cloud
- Redis Client APIs: <https://redis.io/docs/latest/develop/clients/>
 - Connect to Redis servers from languages like JS (NodeJS), Python, Go, Rust, Java, etc.

Tour of the code

Questions? Comments?